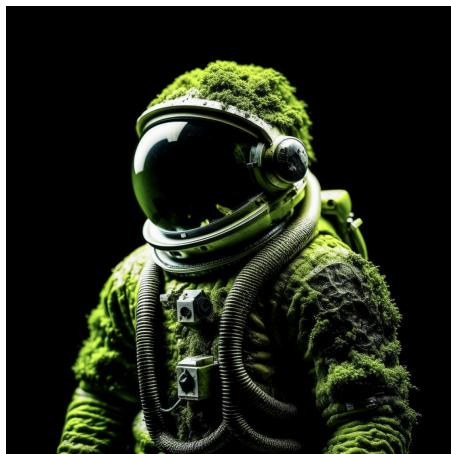


# Kandinsky 2.2

Kandinsky 2.2 is an open source Gen AI multilingual text-to-image [Latent Diffusion Model\(LDM\)](#). In this API, when a text prompt— usually a phrase— is provided to the AI model, an image is generated according to the prompt. For example, if the input text prompt is “A moss covered astronaut with a black background” then the output is an image with an astronaut covered in moss.



**Note:** This API reference document is a sample for one out of four modules . For more information on remaining modules refer to [ai-forever/kandinsky-2.2](#) github repository.

## Authentication

Kandinsky uses a REPLICATE access token to authorize your application. Create a REPLICATE account to get an API token. For more information refer to [REPLICATE website](#)

### curl setup

Set the REPLICATE environment

```
export REPLICATE_API_TOKEN=<paste-your-token-here>
```

```
curl --silent --show-error https://api.replicate.com/v1/predictions \
--request POST \
--header "Authorization: Bearer $REPLICATE_API_TOKEN" \
--header "Content-Type: application/json" \
--header "Prefer: wait" \
```

```
--data @- <<-EOM
{
    "version": "ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb00d2362619f194a",
    "input": {
        "width": 1024,
        "height": 1024
    }
}
EOM
```

## Python setup

Install REPLICATE's python library

```
pip install replicate
```

```
import replicate

input = {
    "width": 1024,
    "height": 1024
}

output = replicate.run(
    "ai-forever/kandinsky-2.2:ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb00d2362619f194a",
    input=input
)
for index, item in enumerate(output):
    with open(f"output_{index}.png", "wb") as file:
        file.write(item.read())
#=> output_0.png written to disk
```

To set up a python environment refer to the [instructions](#).

For more information, refer to [authentication](#) to use Kandinsky.

# API Reference

Kandinsky primarily has four models:

- prior
- text2image
- image fuse
- inpainting

## text2image

The text2image API creates images from the textual prompts provided by the user. It is a decoding diffusion model that converts text embeddings(which are created using an encoder) into image embeddings. These embeddings are later deployed into other Kandinsky's image models such as "Image-to-Image", "Inpainting" and "ControlNet-depth" to modify and increase the accuracy of the image predictions and output.

Run the model

Base URL

```
https://api.replicate.com/v1/predictions
```

curl

```
export REPLICATE_API_TOKEN=<paste-your-token-here>
curl -s -X POST \
-H "Authorization: Bearer $REPLICATE_API_TOKEN" \
-H "Content-Type: application/json" \
-H "Prefer: wait" \
-d $'{
  "version": "ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb00d2362619f194a",
  "input": {
    "width": 1024,
    "height": 1024,
    "prompt": "A moss covered astronaut with a black background",
    "num_outputs": 1,
```

```

    "output_format": "webp",
    "num_inference_steps": 75,
    "num_inference_steps_prior": 25
}
}' \
https://api.replicate.com/v1/predictions

```

## Python

```

pip install diffusers transformers accelerate

from diffusers import AutoPipelineForText2Image
import torch

pipe =
AutoPipelineForText2Image.from_pretrained("kandinsky-community/kandinsky-2-
2-decoder", torch_dtype=torch.float16)
pipe = pipe.to("cuda")

prompt = "A moss covered astronaut with a black background"
negative_prompt = "low quality, bad quality"

image = pipe(prompt=prompt, negative_prompt=negative_prompt,
prior_guidance_scale =1.0, height=768, width=768).images[0]
image.save("portrait.png")

```

**Note:** Refer to [Kandinsky 2.2 repository](#) for the list of modules and [Kandinsky's community](#) for all the other module implementations.

## Parameter list

Input	Type	Parameter	Description
	enum	width	Required width of the output image. Example: 1024

	enum	Height	Required height of the output image. Example: 1024
	string	Prompt	An input string. Example: "A moss covered astronaut with a black background"
	int	Num_outputs	Number of images in the output
	enum	Output_format	The output image format. Example: "jpeg", "png"
	int	Num_inference_steps	The number of denoising or reducing the distortion steps.
	int	num_inference_steps_prior	The number of denoising steps for prior model.

## REST Resources

<code>predictions.create</code>	Create a prediction and get the output.
<code>predictions.get</code>	Get a prediction
<code>predictions.cancel</code>	Cancel a prediction
<code>predictions.list</code>	Lists the predictions

## Request

`predictions.create`

POST/predictions

## curl

```
curl --silent --show-error https://api.replicate.com/v1/predictions \
--request POST \
--header "Authorization: Bearer $REPLICATE_API_TOKEN" \
--header "Content-Type: application/json" \
--header "Prefer: wait" \
--data @- <<-EOM
{
    "version": "ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb00d2362619f194a",
    "input": {
        "width": 1024,
        "height": 1024
    }
}
EOM
```

## Python

```
import replicate

input = {
    "width": 1024,
    "height": 1024
}

output = replicate.run(
    "ai-forever/kandinsky-2.2:ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb
00d2362619f194a",
    input=input
)
for index, item in enumerate(output):
    with open(f"output_{index}.png", "wb") as file:
        file.write(item.read())
#=> output_0.png written to disk
```

## Node.js

```
import { writeFile } from "fs/promises";
import Replicate from "replicate";
const replicate = new Replicate();

const input = {
  width: 1024,
  height: 1024
};

const output = await
replicate.run("ai-forever/kandinsky-2.2:ad9d7879fbffa2874e1d909d1d37d9bc682
889cc65b31f7bb00d2362619f194a", { input });
for (const [index, item] of Object.entries(output)) {
  await writeFile(`output_${index}.png`, item);
}
//=> output_0.png written to disk
```

## predictions.get

GET/predictions/{prediction\_id}

### Parameters

Name	Description
prediction_id	The ID of the prediction to retrieve.

### curl

```
curl --silent https://api.replicate.com/v1/predictions/{prediction_id} \
--request GET \
--header "Authorization: Bearer $REPLICATE_API_TOKEN"
```

Python

```
import replicate

prediction = replicate.predictions.get(prediction_id)
print(prediction)
#=> Prediction(id="xyz...", status="successful", ... )
```

## **predictions.cancel**

```
POST/predictions/{prediction_id}/cancel
```

curl

```
curl --silent
https://api.replicate.com/v1/predictions/{prediction_id}/cancel \
--request POST \
--header "Authorization: Bearer $REPLICATE_API_TOKEN"
```

Python

```
import replicate

prediction = replicate.predictions.get(prediction_id)
prediction.cancel()
print(prediction)
#=> Prediction(id="xyz...", status="canceled", ... )
```

## **Sample Request**

Python

```
output = replicate.run(
```

```
"ai-forever/kandinsky-2.2:ad9d7879fbffa2874e1d909d1d37d9bc682889cc65b31f7bb
00d2362619f194a",
    input={
        "width": 1024,
        "height": 1024,
        "prompt": "A moss covered astronaut with a black background",
        "num_outputs": 1,
        "output_format": "webp",
        "num_inference_steps": 75,
        "num_inference_steps_prior": 25
    }
)
print(output)
```

## Sample Response

JSON

```
{
    "completed_at": "2024-09-12T22:23:00.041096Z",
    "created_at": "2024-09-12T22:22:51.044455Z",
    "data_removed": false,
    "error": null,
    "id": "3tsx2ijbtsa47m4mwbifyiglxm",
    "input": {
        "width": 1024,
        "height": 1024,
        "prompt": "A moss covered astronaut with a black background",
        "num_outputs": 1,
        "num_inference_steps": 75
    },
    "logs": "Using seed: 4697\n  0%          | 0/25 [00:00<?, ?it/s]\n20%|███████| 5/25 [00:00<00:00, 40.13it/s]\n40%|██████████| 10/25\n[00:00<00:00, 40.19it/s]\n60%|██████████| 15/25 [00:00<00:00,\n40.05it/s]\n80%|██████████| 20/25 [00:00<00:00, 40.05it/s]\n100%|██████████| 25/25 [00:00<00:00, 40.05it/s]\n",
    "metrics": {
        "predict_time": 9.029294,
        "total_time": 8.996641
    },
}
```

```
"output": [  
    "https://replicate.delivery/pbxt/Lca3IEjcKoJBBVS6ajR0kK37sDzPsmjYxIcFzxPZp6  
    5wZzTE/out-0.png"  
,  
    "started_at": "2023-07-12T22:22:51.011802Z",  
    "status": "succeeded",  
    "urls": {  
        "get":  
            "https://api.replicate.com/v1/predictions/3tsx2ijbtsha47m4mwbifyiglxm",  
        "cancel":  
            "https://api.replicate.com/v1/predictions/3tsx2ijbtsha47m4mwbifyiglxm/cancel"  
    },  
    "version":  
        "424befb1eae6af8363edb846ae98a11111a39740988baebd279d73fe3ecc92c2"  
}
```

## References:

1. Replicate:: <https://replicate.com/ai-forever/kandinsky-2.2>
2. Github: <https://github.com/ai-forever/Kandinsky-2/blob/main/README.md>
3. Huggingface:  
<https://huggingface.co/kandinsky-community/kandinsky-2-2-prior#text-to-image>
4. Licence: [Apache Licence 2.0](#)